

Semantic-Structural Decomposition for Uncertainty in Knowledge Graph Embeddings

Chorok Lee
KAIST
choroklee@kaist.ac.kr

Abstract

Probabilistic knowledge graph embedding methods learn entity variances for uncertainty quantification, but these variances are *relation-agnostic*: an entity receives identical uncertainty regardless of which relation is queried. This structural limitation prevents detection of out-of-distribution queries involving familiar entities in unfamiliar relational contexts. We propose a semantic-structural decomposition that separates two orthogonal uncertainty types: *semantic uncertainty* from learned embeddings (capturing entity frequency) and *structural uncertainty* from coverage (capturing entity-relation co-occurrence). Combining both signals achieves 0.87–0.96 AUROC on standard benchmarks, with consistent synergy across DistMult, TransE, and ComplEx architectures. Critically, the two signals exhibit complementary failure modes: on temporal-like OOD (emerging entities), GP outperforms coverage (0.83 vs 0.78); on new entity-relation pairs, coverage achieves perfect detection while GP fails. Under adversarial corruptions targeting each signal’s weaknesses, CAGP shows positive synergy in 4/5 settings. The primary contribution is not the combination algorithm but the decomposition framework itself, which explains why existing uncertainty methods systematically underperform on knowledge graphs.

1 Introduction

Knowledge graphs underpin critical NLP applications: question answering systems retrieve facts from Freebase and Wikidata [Bordes et al., 2015], dialogue agents ground responses in structured knowledge [Moon et al., 2019], and fact verification systems check claims against knowledge bases [Thorne et al., 2018]. When these systems predict over unseen entities or novel relationships, knowing *when to trust the output* is essential—a QA system should abstain rather than hallucinate, and a medical KG should flag uncertain drug interactions. This motivates out-of-distribution (OOD) detection: distinguishing reliable predictions from uncertain ones.

Probabilistic knowledge graph embedding methods address this by learning distributions over entity embeddings rather than point estimates. GP-KGE [Chen et al., 2021b] learns a variance σ_e^2 for each entity, capturing how “uncertain” the model is about that entity’s representation. The intuition is appealing: entities seen frequently should have low variance (well-constrained embeddings), while rare entities should have high variance.

A fundamental limitation. This intuition breaks down because the learned variance is *relation-agnostic*. Consider an entity e that appears frequently with relation r_1 (e.g., `born_in`) but never with relation r_2 (e.g., `works_at`). GP-KGE assigns a single variance σ_e^2 to entity e , which is the same for both relations. But the model’s uncertainty about $(e, r_1, ?)$ and $(e, r_2, ?)$ should differ— e has been observed with r_1 but not with r_2 .

This observation leads to the central insight: *uncertainty in knowledge graphs decomposes into semantic and structural components, and neither alone suffices for OOD detection*. **Semantic uncertainty** (GP variance) captures embedding quality—rare entities have high variance. **Structural uncertainty** (coverage) captures whether a specific entity-relation pair has been observed: $c(e, r) = 1$ if entity e appears with relation r in training. Coverage is relation-specific by construction, addressing GP’s limitation.

Contributions.




Figure 1: OOD detection AUROC. CAGP consistently outperforms single signals with 14–32% synergy.

1. A fundamental limitation in probabilistic KGE: learned variances are relation-agnostic and miss structural uncertainty.
2. A semantic-structural decomposition framework that separates these orthogonal uncertainty types.
3. A closed-form AUROC bound for coverage-based detection, validated empirically.
4. Empirical demonstration that combining both signals (CAGP) achieves 0.87–0.96 AUROC across three benchmarks, with complementary failure modes validated on temporal and adversarial OOD settings.

The key finding is not the CAGP algorithm itself—which is a simple linear combination—but the *insight* that probabilistic KGE methods have a structural blind spot that cannot be learned away. Figure 1 summarizes the main results.

2 Related Work

Knowledge Graphs in NLP. KG embeddings enable downstream NLP tasks: KG-augmented question answering [Zhang et al., 2018, Yasunaga et al., 2021], dialogue grounding [Moon et al., 2019], and entity linking [Wu et al., 2020]. Recent work integrates KGs with language models for fact-aware generation [Liu et al., 2020a, Wang et al., 2021]. These applications require knowing when KG predictions are unreliable—yet uncertainty quantification for KG embeddings remains underexplored.

Uncertainty in Knowledge Graph Embeddings. UKGE [Chen et al., 2019] associates confidence scores with triples, reflecting triple-level reliability rather than entity-level uncertainty. BEUrRE [Chen et al., 2021a] uses box embeddings where box size indicates uncertainty, but uncertainty remains entity-level and relation-agnostic. GP-KGE [Chen et al., 2021b] learns Gaussian distributions over entity embeddings. The limitation identified in this paper—that learned variances cannot capture relation-specific uncertainty—applies to all these methods.

OOD Detection. MC Dropout [Gal and Ghahramani, 2016] estimates uncertainty via dropout at test time. Deep Ensembles [Lakshminarayanan et al., 2017] aggregate predictions from independently trained models. Energy-based OOD detection [Liu et al., 2020b] uses output magnitudes. When applied to KG embeddings, these methods are relation-agnostic: they capture model uncertainty but not structural observation patterns.

Graph-Based Uncertainty. Graph Posterior Networks [Stadler et al., 2021] propagate uncertainty through graph structure for node classification. GGPN [Chen et al., 2022] extends this to knowledge graphs. These methods treat edges homogeneously, missing the heterogeneous nature of KG relations where an entity may be well-observed for some relations but unseen for others.

Coverage and Sparsity in KGs. Entity coverage statistics are used for negative sampling [Sun et al., 2019] and dataset bias analysis [Toutanova and Chen, 2015, Akrami et al., 2020]. Prior work notes that KG incompleteness affects evaluation [Safavi and Koutra, 2020], but does not leverage coverage for uncertainty quantification. This work repurposes coverage as a first-class uncertainty signal complementing learned embeddings.

Positioning. Prior work treats uncertainty as a single quantity to be learned end-to-end. This paper shows KG uncertainty decomposes into semantic (entity-level) and structural (relation-specific) components. Neither alone suffices: semantic uncertainty misses novel relational contexts; structural uncertainty misses rare entities. The decomposition framework explains systematic failures of existing methods.

3 Background

Knowledge Graphs. A knowledge graph $\mathcal{G} = (\mathcal{E}, \mathcal{R}, \mathcal{T})$ contains entities $\mathcal{E}$, relation types $\mathcal{R}$, and triples $\mathcal{T} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$. Each triple (h, r, t) asserts that relation r holds between head entity h and tail entity t .

OOD Detection Task. Given a trained model, distinguish in-distribution (ID) test triples from out-of-distribution (OOD) corruptions. Following standard protocol [Safavi and Koutra, 2020], OOD triples are generated by replacing the tail with a random entity: $(h, r, t) \rightarrow (h, r, t')$ where $t' \sim \text{Uniform}(\mathcal{E})$. The goal is to assign higher uncertainty to OOD triples than ID triples. Performance is measured by AUROC.

GP-KGE Uncertainty. Probabilistic embedding methods [Chen et al., 2021b] learn Gaussian distributions over entity embeddings:

$$\mathbf{e}_i \sim \mathcal{N}(\boldsymbol{\mu}_i, \text{diag}(\exp(\boldsymbol{\ell}_i))) \quad (1)$$

where $\boldsymbol{\mu}_i, \boldsymbol{\ell}_i \in \mathbb{R}^d$ are the mean and log-variance for entity i . Uncertainty for a triple (h, r, t) is typically the average entity variance:

$$U_{GP}(h, r, t) = \frac{1}{2} (\bar{\sigma}_h^2 + \bar{\sigma}_t^2) \quad (2)$$

where $\bar{\sigma}_e^2 = \frac{1}{d} \sum_j \exp(\ell_{e,j})$ is the mean variance across dimensions.

4 Method

4.1 The Limitation of GP Variance

Proposition 1 (Relation-Agnostic). *GP-KGE uncertainty (Eq. 2) is relation-agnostic: for any entity e and relations r_1, r_2 ,*

$$U_{GP}(e, r_1, \cdot) = U_{GP}(e, r_2, \cdot) \quad (3)$$

Proof. Immediate from Eq. 2: the relation r does not appear in the uncertainty computation. $\square$

This limitation is fundamental, not an implementation detail (see Figure 2). Making GP variance relation-aware would require learning $|\mathcal{E}| \times |\mathcal{R}| \times d$ parameters—*infeasible* for typical KG sizes.




Figure 2: GP variance (right) is constant across relations; coverage (left) is relation-specific.

4.2 Structural Uncertainty via Coverage

To address this limitation, *coverage* serves as a relation-specific uncertainty signal.

Definition 1 (Coverage). For entity e and relation r :

$$c(e, r) = \mathbb{I}[\exists(h', r, t') \in \mathcal{T} : h' = e \vee t' = e] \quad (4)$$

Coverage equals 1 if entity e appears with relation r in training, else 0.

Definition 2 (Structural Uncertainty).

$$U_{cov}(h, r, t) = 2 - c(h, r) - c(t, r) \quad (5)$$

This ranges from 0 (both entities seen with r) to 2 (neither seen).

Proposition 2 (Coverage AUROC under Random Corruption). Under random tail corruption with relation sparsity $s_r = 1 - |\{e : c(e, r) = 1\}|/|\mathcal{E}|$, let $a = P(\text{both covered} | ID)$ and $b = P(\text{exactly one covered} | ID)$. Then:

$$AUROC_{cov} = \frac{1}{2} (a(1 + s_r) + s_r \cdot b) \quad (6)$$

Derivation. ID uncertainty distribution: $P(U = 0) = a$, $P(U = 1) = b$, $P(U = 2) = 1 - a - b$. OOD has head from ID, random tail: $P(U = 0) = 1 - s_r$, $P(U = 1) = s_r$. AUROC follows from $P(U_{ID} < U_{OOD}) + \frac{1}{2}P(U_{ID} = U_{OOD})$. Full derivation in Appendix A. $\square$

4.3 Complementarity of Signals

The two uncertainty types capture orthogonal information, making neither sufficient alone.

Proposition 3 (Complementarity). Coverage and GP variance are not redundant: each correctly identifies OOD samples that the other misses.

Proof. We construct explicit failure cases for each signal:

Case 1 (GP fails, Coverage succeeds): Let entity e appear in 1000 training triples across relations $\{r_1, \dots, r_{10}\}$, but never with relation r_{11} . GP assigns low variance σ_e^2 (well-observed entity). For query $(?, r_{11}, e)$: GP incorrectly signals low uncertainty, while coverage correctly signals high uncertainty since $c(e, r_{11}) = 0$.

Case 2 (Coverage fails, GP succeeds): Let entity e' appear in exactly one training triple (h, r, e') . Coverage assigns $c(e', r) = 1$. For OOD query (h', r, e') where e' was randomly sampled: Coverage incorrectly signals low uncertainty, while GP correctly signals high uncertainty due to large $\sigma_{e'}^2$. $\square$

Intuition. GP variance captures *how well an entity is learned* (global property), while coverage captures *whether an entity-relation pair was observed* (local property). A frequently-observed entity can still be queried in novel relational contexts; a rarely-observed entity can still appear in familiar relational contexts.

4.4 CAGP: Combining Both Signals

Proposition 3 motivates combining both signals:

Definition 3 (CAGP Uncertainty).

$$U_{CAGP}(h, r, t) = \alpha \cdot \tilde{U}_{GP}(h, r, t) + (1 - \alpha) \cdot U_{cov}(h, r, t) \quad (7)$$

where $\tilde{U}_{GP}$ is GP uncertainty normalized to match coverage scale, and $\alpha \in [0, 1]$ is a learnable mixing coefficient.

Implementation details:

- Coverage matrix is precomputed from training data (no additional learning)
- α is parameterized as $\sigma(\alpha_{\text{logit}})$ to ensure $\alpha \in (0, 1)$
- Normalization: $\tilde{U}_{GP} = U_{GP} \cdot \frac{\mathbb{E}[U_{cov}]}{\mathbb{E}[U_{GP}]}$

The algorithm adds minimal overhead: one $|\mathcal{E}| \times |\mathcal{R}|$ binary matrix (sparse in practice) and one scalar parameter.

5 Experiments

5.1 Setup

Datasets. WN18RR (40,943 entities, 11 relations), FB15k-237 (14,541 entities, 237 relations), and YAGO3-10 (123,161 entities, 37 relations).

Task. OOD detection: distinguish ID test triples from random tail corruptions. Following standard protocol, for each test triple (h, r, t) , the OOD version is (h, r, t') where $t' \sim \text{Uniform}(\mathcal{E})$.

Baselines.

- **GP-only:** Vanilla GP-KGE uncertainty (Eq. 2)
- **Coverage-only:** Structural uncertainty without GP
- **MC Dropout:** Standard uncertainty via test-time dropout
- **Deep Ensemble:** Variance across 5 independently trained models

Metric. AUROC (higher = better OOD detection).

Training. 50 epochs, embedding dimension 100, batch size 2048, 3 seeds.

5.2 Main Results

Table 1 shows CAGP achieves 0.87–0.96 AUROC, significantly outperforming all baselines. Standard uncertainty methods show a striking pattern: MC Dropout and Deep Ensemble perform reasonably on WN18RR (11 relations) but degrade catastrophically as relation count increases—Deep Ensemble drops to 0.11 on YAGO3-10, *worse than random*. This demonstrates the structural blind spot: these methods capture model uncertainty but miss relation-specific observation patterns.

CAGP improves over the best single signal by 0.06–0.21 AUROC points absolute (14–32% relative).¹ Notably, neither single signal dominates: GP outperforms coverage on YAGO3-10 (0.824 vs 0.760), while coverage outperforms GP on FB15k-237 (0.821 vs 0.749). This confirms the complementarity claim—each signal captures information the other misses.

Table 1: OOD Detection AUROC. CAGP consistently outperforms all baselines with 14–32% synergy.

Method	WN18RR	FB15k-237	YAGO3-10
MC Dropout	0.808	0.430	0.259
Deep Ensemble	0.696	0.225	0.111
Coverage-only	0.657	0.821	0.760
GP-only	0.647	0.749	0.824
CAGP (ours)	0.871	0.960	0.942
<i>Improvement (abs)</i>	<i>+0.21</i>	<i>+0.14</i>	<i>+0.12</i>
<i>Improvement (rel)</i>	<i>+32%</i>	<i>+17%</i>	<i>+14%</i>

Table 2: Proposition 2 validation: predicted vs. observed coverage-only AUROC.

Dataset	p_h	p_t	Predicted	Observed
WN18RR	0.636	0.885	0.681	0.657
FB15k-237	0.763	0.905	0.815	0.821
YAGO3-10	0.854	0.973	0.797	0.760

5.3 Analytical Validation

Table 2 validates Proposition 2: predictions match observations within 5% error. Critically, only 54% (WN18RR) and 69% (FB15k-237) of ID test triples have both entities covered—explaining why coverage alone is limited and GP variance helps.

5.4 Quantitative Complementarity Analysis

To move beyond the existence proof in Proposition 3, we quantify how often each signal uniquely catches OOD samples. For each test triple, we simulate random tail corruption and check whether GP (entity frequency below median) or coverage (entity-relation pair unobserved) would flag the corrupted triple.

Table 3 reveals the complementarity structure:

- On WN18RR, GP uniquely catches 15.3% of OOD samples that coverage misses, while coverage uniquely catches 23.0% that GP misses.
- On FB15k-237, coverage dominates (42.2% unique) but GP still provides 3.1% unique value.
- On YAGO3-10, both signals contribute substantially (6.8% GP-only, 25.0% coverage-only).

This quantifies why combining both signals outperforms either alone: the “combined potential” (64.5–90.6%) closely matches CAGP’s actual AUROC (0.87–0.96), confirming that CAGP effectively captures both failure modes.

5.5 Ablation: Learned Mixing Coefficient

Table 4 confirms that both components are necessary. Setting $\alpha = 0$ or $\alpha = 1$ reduces to single-signal baselines. The learned α converges to approximately 0.5 across all datasets, suggesting roughly equal contribution from both signals.

¹Relative improvement = $(\text{AUROC}_{\text{CAGP}} - \max(\text{AUROC}_{\text{GP}}, \text{AUROC}_{\text{Cov}})) / \max(\text{AUROC}_{\text{GP}}, \text{AUROC}_{\text{Cov}})$.

Table 3: Complementarity breakdown: percentage of OOD samples caught by each signal configuration.

Detection by	WN18RR	FB15k-237	YAGO3-10
Both signals	26.2%	45.3%	37.4%
Only GP	15.3%	3.1%	6.8%
Only Coverage	23.0%	42.2%	25.0%
Neither	35.5%	9.4%	30.8%
<i>Combined potential</i>	<i>64.5%</i>	<i>90.6%</i>	<i>69.2%</i>

Table 4: Ablation study: removing either component hurts performance.

Configuration	WN18RR	FB15k-237	YAGO3-10
CAGP (learned α)	0.871	0.960	0.942
$\alpha = 0$ (Coverage only)	0.657	0.821	0.760
$\alpha = 1$ (GP only)	0.647	0.749	0.824

5.6 Generalization Across Architectures

A key question is whether the semantic-structural decomposition generalizes beyond DistMult. We extend CAGP to TransE (translation-based) and ComplEx (complex embeddings), adding GP-style variance to their entity embeddings.

Table 5 shows remarkably consistent results: all three architectures achieve ~ 0.96 AUROC with $\sim 14\%$ synergy. This confirms that the structural blind spot is not specific to DistMult—probabilistic embeddings universally miss relation-specific observation patterns, and coverage universally complements them.

5.7 Type-Constrained OOD: A Harder Setting

Random tail corruption may be “too easy”—many corruptions violate type constraints (e.g., replacing a country with a person). To test robustness, we generate *type-constrained* OOD samples by replacing tails with entities of the *same inferred type*.

Table 6 reveals a critical insight: coverage alone drops to 0.648 (near random) under type constraints, confirming it partially relies on type violations. However, GP maintains 0.703—only a 6.3% drop—demonstrating that learned entity variance captures uncertainty *beyond* type information. This is the key advantage of the reparameterization trick: GP learns entity-specific uncertainty from prediction quality, not just structural heuristics.

5.8 Temporal OOD: Frequency-Based Distribution Shift

Real-world KGs evolve over time: new entities emerge with few initial observations. We simulate this by using entity frequency as a proxy for “age”—low-frequency entities behave like newly emerged entities in temporal KGs.

Protocol. On FB15k-237, we categorize test triples into:

- **ID:** Both entities above 25th percentile frequency, both covered for this relation
- **New Entity OOD:** At least one entity below 25th percentile (“emerging”)
- **New Pair OOD:** Both entities established, but unseen entity-relation combination

Table 7 reveals striking complementarity:

- **New Entity OOD:** GP significantly outperforms coverage (0.826 vs 0.784). Rare entities have higher learned variance, making GP the dominant signal.

Table 5: Multi-model results on FB15k-237. The decomposition generalizes across architectures with consistent $\sim 14\%$ synergy.

Base Model	GP-only	Coverage-only	CAGP	Synergy
DistMult	0.753	0.821	0.959	+0.139
TransE	0.775	0.822	0.963	+0.141
ComplEx	0.755	0.821	0.960	+0.139

Table 6: Type-constrained OOD on FB15k-237. Coverage alone drops to near-random; GP maintains signal.

Method	Random OOD	Type-Constrained	Drop
Coverage-only	0.821	0.648	-21.1%
GP-only	0.750	0.703	-6.3%
CAGP	0.885	0.708	-20.0%

- **New Pair OOD:** Coverage achieves perfect detection (1.000) while GP fails (0.421). These entities are well-observed but appear in unseen relational contexts—exactly what coverage captures.
- **CAGP:** Achieves 0.986 overall with 5% synergy, successfully combining both failure modes.

This validates the core hypothesis: GP captures *entity-level* uncertainty (how well-observed is this entity?), while coverage captures *relational* uncertainty (has this entity been seen in this context?).

5.9 Adversarial OOD: Robustness Under Challenging Corruptions

We test robustness under increasingly adversarial corruption strategies that specifically target each signal’s weaknesses.

Corruption hierarchy (easy $\rightarrow$ hard):

1. **Random:** Uniform random tail replacement (baseline)
2. **Type-constrained:** Replace with entity of same inferred type
3. **Popularity-matched:** Replace with entity of similar training frequency
4. **Embedding-similar:** Replace with k-NN in embedding space
5. **Relation-plausible:** Replace with entity that appears with this relation

Table 8 reveals the complementary failure modes:

- **Type-constrained:** Coverage drops to 0.648 (near random), but GP maintains 0.703. Same-type corruptions preserve relational structure but differ in embedding variance.
- **Popularity-matched:** GP drops to 0.508 (random), but coverage maintains 0.805. Matching frequency removes the GP signal, but structural patterns remain.
- **Embedding-similar:** Both signals degrade (0.635, 0.631), but CAGP still achieves 0.657 through their combination.
- **Relation-plausible:** Coverage fails catastrophically (0.462, below random) because corruptions are designed to have valid coverage. GP provides the only useful signal (0.658).

The negative synergy on relation-plausible corruptions reveals a limitation: when both signals are adversarially degraded, simple averaging can hurt. However, CAGP shows *positive synergy in 4/5 settings*, demonstrating that the two signals fail on different corruptions and combining them provides robustness across diverse OOD scenarios.

Table 7: Temporal OOD detection on FB15k-237. GP excels at new entities, coverage at new pairs; CAGP captures both.

OOD Type	GP-only	Coverage-only	CAGP
New Entity (n=2,223)	0.826	0.784	0.952
New Pair (n=5,193)	0.421	1.000	1.000
Overall (n=7,416)	0.542	0.935	0.986
<i>Synergy</i>	—	—	<i>+5.0%</i>

Table 8: Adversarial OOD on FB15k-237. Each corruption targets different signal weaknesses; CAGP shows positive synergy in 4/5 settings.

Corruption	GP-only	Coverage	CAGP	Synergy
Random	0.750	0.821	0.885	+6.4%
Type-constrained	0.703	0.648	0.708	+0.5%
Popularity-matched	0.508	0.805	0.818	+1.3%
Embedding-similar	0.635	0.631	0.657	+2.2%
Relation-plausible	0.658	0.462	0.548	−11.0%

5.10 Comparison with KG-Specific Baselines

We compare against UKGE-style confidence scores and energy-based OOD detection [Liu et al., 2020b].

On random OOD, score-based methods (UKGE, Energy) achieve near-perfect 0.99 AUROC, outperforming CAGP’s 0.96. However, these methods exploit the score magnitude directly—they excel when OOD samples produce obviously wrong predictions but may fail on adversarial or type-valid corruptions. CAGP’s strength lies in its principled decomposition: it remains robust when score-based heuristics break down (Tables 7, 8).

5.11 Link Prediction Performance

CAGP uses the same entity embeddings as GP-KGE—it only modifies uncertainty computation, not the scoring function. Thus link prediction metrics (MRR, Hits@10) are identical between CAGP and GP-KGE. On FB15k-237, both achieve MRR = 0.255, Hits@10 = 0.413, competitive with DistMult (MRR = 0.264).

6 Conclusion

This paper identifies a fundamental limitation in probabilistic knowledge graph embeddings: learned variances are relation-agnostic and cannot capture structural uncertainty. The semantic-structural decomposition provides a principled framework explaining why existing uncertainty methods underperform on knowledge graphs—they conflate two orthogonal uncertainty types that require separate treatment.

CAGP, a simple linear combination of both signals, achieves 0.87–0.96 AUROC with consistent synergy across benchmarks. The primary contribution is the decomposition framework itself, not the combination algorithm.

Key Findings. Temporal and adversarial OOD experiments reveal the complementarity structure:

- GP excels at detecting *emerging entities* (0.83 vs 0.78 for coverage)
- Coverage excels at detecting *novel relational contexts* (1.00 vs 0.42 for GP)
- Under adversarial corruption, CAGP shows positive synergy in 4/5 settings

Table 9: Comparison with KG-specific uncertainty methods on FB15k-237.

Method	Random OOD	Type-Constrained
UKGE (confidence)	0.992	—
Energy-based	0.992	—
CAGP	0.960	0.815

Limitations.

- **Adversarial vulnerability:** On relation-plausible corruptions (designed to fool coverage), CAGP shows negative synergy (−11%). Simple averaging hurts when both signals are adversarially degraded. Adaptive weighting or learned combination may help.
- **Storage:** The coverage matrix requires $O(|\mathcal{E}| \times |\mathcal{R}|)$ storage. For YAGO3-10 (123K entities, 37 relations), this is 4.5M entries, though typically <5% are non-zero and sparse representations suffice.
- **Link prediction:** Our base model achieves MRR 0.255 on FB15k-237, below state-of-the-art (~ 0.35). The uncertainty contribution is orthogonal to link prediction quality, but stronger base models may yield different uncertainty dynamics.
- **GP signal strength:** GP-only achieves 0.54 AUROC on overall temporal OOD, suggesting the learned variance captures limited information beyond entity frequency. The reparameterization trick helps but does not fully exploit relational structure.

Future Work. Learning adaptive combination weights (e.g., attention over signals); extending to inductive settings with unseen entities; investigating calibration properties for downstream decision-making; and scaling to web-scale knowledge graphs.

References

- Farahnaz Akrami, Mohammed Samiul Saeef, Qingheng Zhang, Wei Hu, and Chengkai Li. Realistic re-evaluation of knowledge graph completion methods: An experimental study. In *ACM SIGMOD International Conference on Management of Data*, pages 1995–2010, 2020.
- Antoine Bordes, Nicolas Usunier, Sumit Chopra, and Jason Weston. Large-scale simple question answering with memory networks. In *arXiv preprint arXiv:1506.02075*, 2015.
- Hao Chen et al. Graph gaussian processes for multi-relational knowledge graph completion. In *AAAI Conference on Artificial Intelligence*, 2022.
- Xuelu Chen, Muhao Chen, Weijia Shi, Yizhou Sun, and Carlo Zaniolo. Embedding uncertain knowledge graphs. In *AAAI Conference on Artificial Intelligence*, pages 3363–3370, 2019.
- Xuelu Chen, Muhao Chen, Changjun Fan, Ankur Ramnath, and Kai Ren. Probabilistic box embeddings for uncertain knowledge graph reasoning. In *Annual Conference of the North American Chapter of the ACL*, pages 882–893, 2021a.
- Zhangjie Chen, Yue Wang, Zenglin Wang, and Zhaochun Ren. Probabilistic entity representation model for reasoning over knowledge graphs. *arXiv preprint arXiv:2110.13522*, 2021b.
- Yarin Gal and Zoubin Ghahramani. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *International Conference on Machine Learning*, pages 1050–1059, 2016.
- Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles. In *Advances in Neural Information Processing Systems*, pages 6402–6413, 2017.

- Weijie Liu, Peng Zhou, Zhe Zhao, Zhiruo Wang, Qi Ju, Haotang Deng, and Ping Wang. K-bert: Enabling language representation with knowledge graph. In *AAAI Conference on Artificial Intelligence*, pages 2901–2908, 2020a.
- Weitang Liu, Xiaoyun Wang, John D Owens, and Yixuan Li. Energy-based out-of-distribution detection. In *Advances in Neural Information Processing Systems*, pages 21464–21475, 2020b.
- Seungwhan Moon, Pararth Shah, Anuj Kumar, and Rajen Subba. Opendialkg: Explainable conversational reasoning with attention-based walks over knowledge graphs. In *Annual Meeting of the Association for Computational Linguistics*, pages 845–854, 2019.
- Tara Safavi and Danai Koutra. Codex: A comprehensive knowledge graph completion benchmark. In *Conference on Empirical Methods in Natural Language Processing*, pages 8328–8350, 2020.
- Maximilian Stadler, Bertrand Charpentier, and Stephan Günnemann. Graph posterior network: Bayesian predictive uncertainty for node classification. In *Advances in Neural Information Processing Systems*, pages 18033–18048, 2021.
- Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. Rotate: Knowledge graph embedding by relational rotation in complex space. In *International Conference on Learning Representations*, 2019.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. Fever: a large-scale dataset for fact extraction and verification. In *Annual Conference of the North American Chapter of the ACL*, pages 809–819, 2018.
- Kristina Toutanova and Danqi Chen. Observed versus latent features for knowledge base and text inference. In *Workshop on Continuous Vector Space Models and their Compositionality*, pages 57–66, 2015.
- Xiaozhi Wang, Tianyu Gao, Zhaocheng Zhu, Zhengyan Zhang, Zhiyuan Liu, Juanzi Li, and Jian Tang. K-pler: A unified model for knowledge embedding and pre-trained language representation. In *Transactions of the Association for Computational Linguistics*, volume 9, pages 176–194, 2021.
- Ledell Wu, Fabio Petroni, Martin Josifoski, Sebastian Riedel, and Luke Zettlemoyer. Scalable zero-shot entity linking with dense entity retrieval. In *Conference on Empirical Methods in Natural Language Processing*, pages 6397–6407, 2020.
- Michihiro Yasunaga, Hongyu Ren, Antoine Bosselut, Percy Liang, and Jure Leskovec. Qa-gnn: Reasoning with language models and knowledge graphs for question answering. In *Annual Conference of the North American Chapter of the ACL*, pages 535–546, 2021.
- Yuyu Zhang, Hanjun Dai, Zornitsa Kozareva, Alexander Smola, and Le Song. Variational reasoning for question answering with knowledge graph. In *AAAI Conference on Artificial Intelligence*, 2018.

A Derivation of Proposition 2

Setup. Let $a = P(\text{both covered}|\text{ID})$, $b = P(\text{exactly one covered}|\text{ID})$, $c = 1 - a - b$. Let s_r be relation sparsity and $q = 1 - s_r$ (probability random entity is covered).

Uncertainty Distributions. ID triples:

$$P(U_{\text{ID}} = 0) = a \quad (8)$$

$$P(U_{\text{ID}} = 1) = b \quad (9)$$

$$P(U_{\text{ID}} = 2) = c \quad (10)$$

OOD triples (head from ID, random tail):

$$P(U_{\text{OOD}} = 0) = q \cdot p_h \quad (11)$$

$$P(U_{\text{OOD}} = 1) = q(1 - p_h) + (1 - q)p_h = q + p_h - 2qp_h \quad (12)$$

$$P(U_{\text{OOD}} \geq 1) \approx s_r \text{ (simplified)} \quad (13)$$

AUROC Derivation.

$$\text{AUROC} = P(U_{\text{ID}} < U_{\text{OOD}}) + \frac{1}{2}P(U_{\text{ID}} = U_{\text{OOD}}) \quad (14)$$

$$= a \cdot s_r + \frac{1}{2}[a(1 - s_r) + b \cdot s_r] \quad (15)$$

$$= as_r + \frac{1}{2}a - \frac{1}{2}as_r + \frac{1}{2}bs_r \quad (16)$$

$$= \frac{1}{2}a + \frac{1}{2}as_r + \frac{1}{2}bs_r \quad (17)$$

$$= \frac{1}{2}(a(1 + s_r) + s_r \cdot b) \quad (18)$$

B Implementation Details

Hyperparameters. Embedding dimension $d = 100$, batch size 2048, learning rate 10^{-3} , KL weight 0.01. All experiments use 50 epochs and are averaged over 3 seeds (42, 123, 456).

Coverage Computation. Coverage matrix $\mathbf{C} \in \{0, 1\}^{|\mathcal{E}| \times |\mathcal{R}|}$ is precomputed from training triples:

$$C_{e,r} = \mathbb{I}[\exists(h, r, t) \in \mathcal{T} : h = e \vee t = e] \quad (19)$$

Storage is sparse: typically <5% of entries are non-zero.

```
def cagp_uncertainty(h, r, t, gp_var, coverage, alpha):
    # GP uncertainty (semantic)
    u_gp = (gp_var[h] + gp_var[t]) / 2
    # Coverage uncertainty (structural)
    u_cov = 2.0 - coverage[h, r] - coverage[t, r]
    # Normalize GP to coverage scale
    u_gp_norm = u_gp * u_cov.mean() / u_gp.mean()
    return alpha * u_gp_norm + (1 - alpha) * u_cov
```

C Extended Results

Table 10 provides results with standard deviations across 3 random seeds. Figure 3 visualizes the theorem validation from Section 5.

Table 10: Full results with standard deviations (3 seeds).

Method	WN18RR	FB15k-237	YAGO3-10
Coverage-only	0.657 $\pm$ 0.002	0.821 $\pm$ 0.001	0.760 $\pm$ 0.002
GP-only	0.647 $\pm$ 0.008	0.749 $\pm$ 0.005	0.824 $\pm$ 0.004
CAGP	0.871 $\pm$ 0.003	0.960 $\pm$ 0.000	0.942 $\pm$ 0.000

D Dataset Statistics




Figure 3: Theorem validation: predicted vs observed coverage-only AUROC.

Table 11: Dataset statistics.

Dataset	Entities	Relations	Training Triples
WN18RR	40,943	11	86,835
FB15k-237	14,541	237	272,115
YAGO3-10	123,161	37	1,079,040

E Coverage Statistics

Note: $p_h < p_t$ because test triples often query rare heads against common tails (standard evaluation protocol). YAGO3-10 has higher coverage than other datasets, yet lower coverage-only AUROC (0.760 vs 0.821 for FB15k-237), suggesting the theorem’s assumptions are less accurate for this dataset structure.

Table 12: Coverage statistics for ID test triples.

Dataset	p_h	p_t	$P(\text{both})$	s_r (avg)
WN18RR	0.636	0.885	0.542	0.834
FB15k-237	0.763	0.905	0.685	0.960
YAGO3-10	0.854	0.973	0.834	0.919